

Streaming Systems for Real-Time Analytics

RJ Nowling

Memphis.dev

MadPy Meetup

August 10, 2023

Introducing Myself



- Developer Advocate at Memphis.dev
- Associate Professor of Computer Science at the Milwaukee School of Engineering (5 years)
- ML engineer in industry (4 years)

BigPetStore



Simulated Transactions

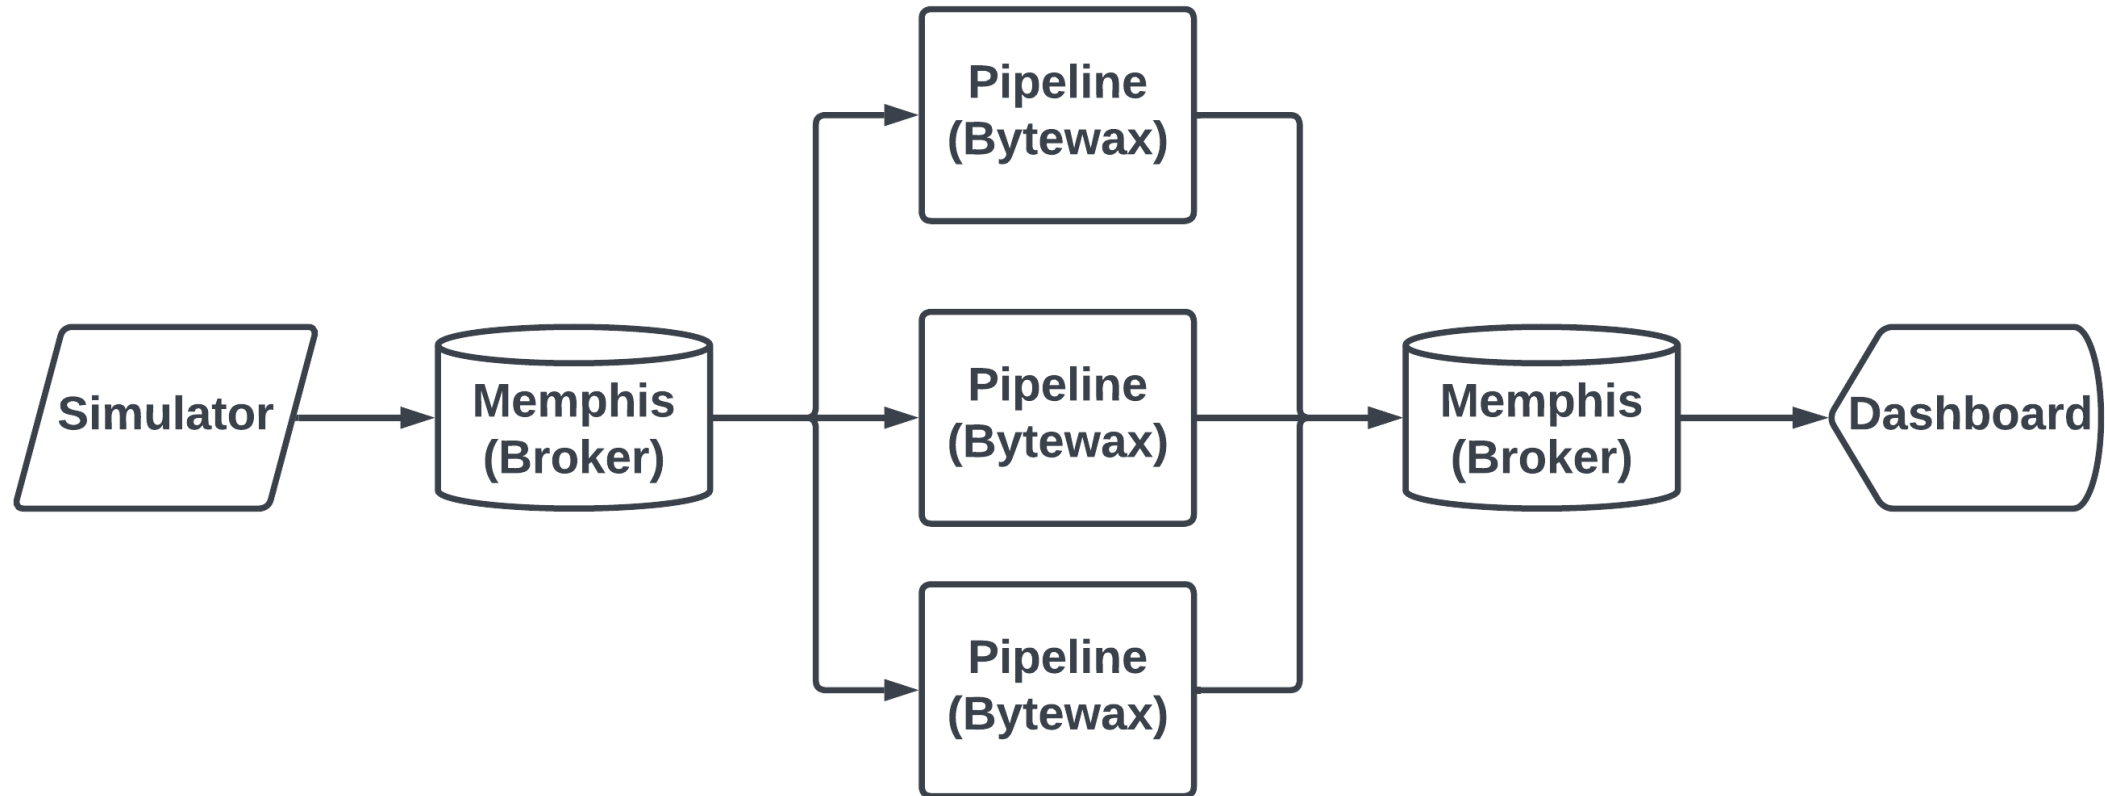
```
{  
  "transaction_id": 1107316,  
  "customer_id": 69770,  
  "timestamp": "2023-01-01T05:59:40.799777",  
  "line_items": [  
    { "item_id": 10, "quantity": 1 },  
    { "item_id": 5, "quantity": 2 }  
  ]  
}
```

1.5 million transactions
1 year of time

Analyses

- Weekly transaction volume (tumbling window)
- Week-over-week change in transaction volume (sliding window)
- Week-over-week change in product sales volume (sliding window, group by key)

Architecture

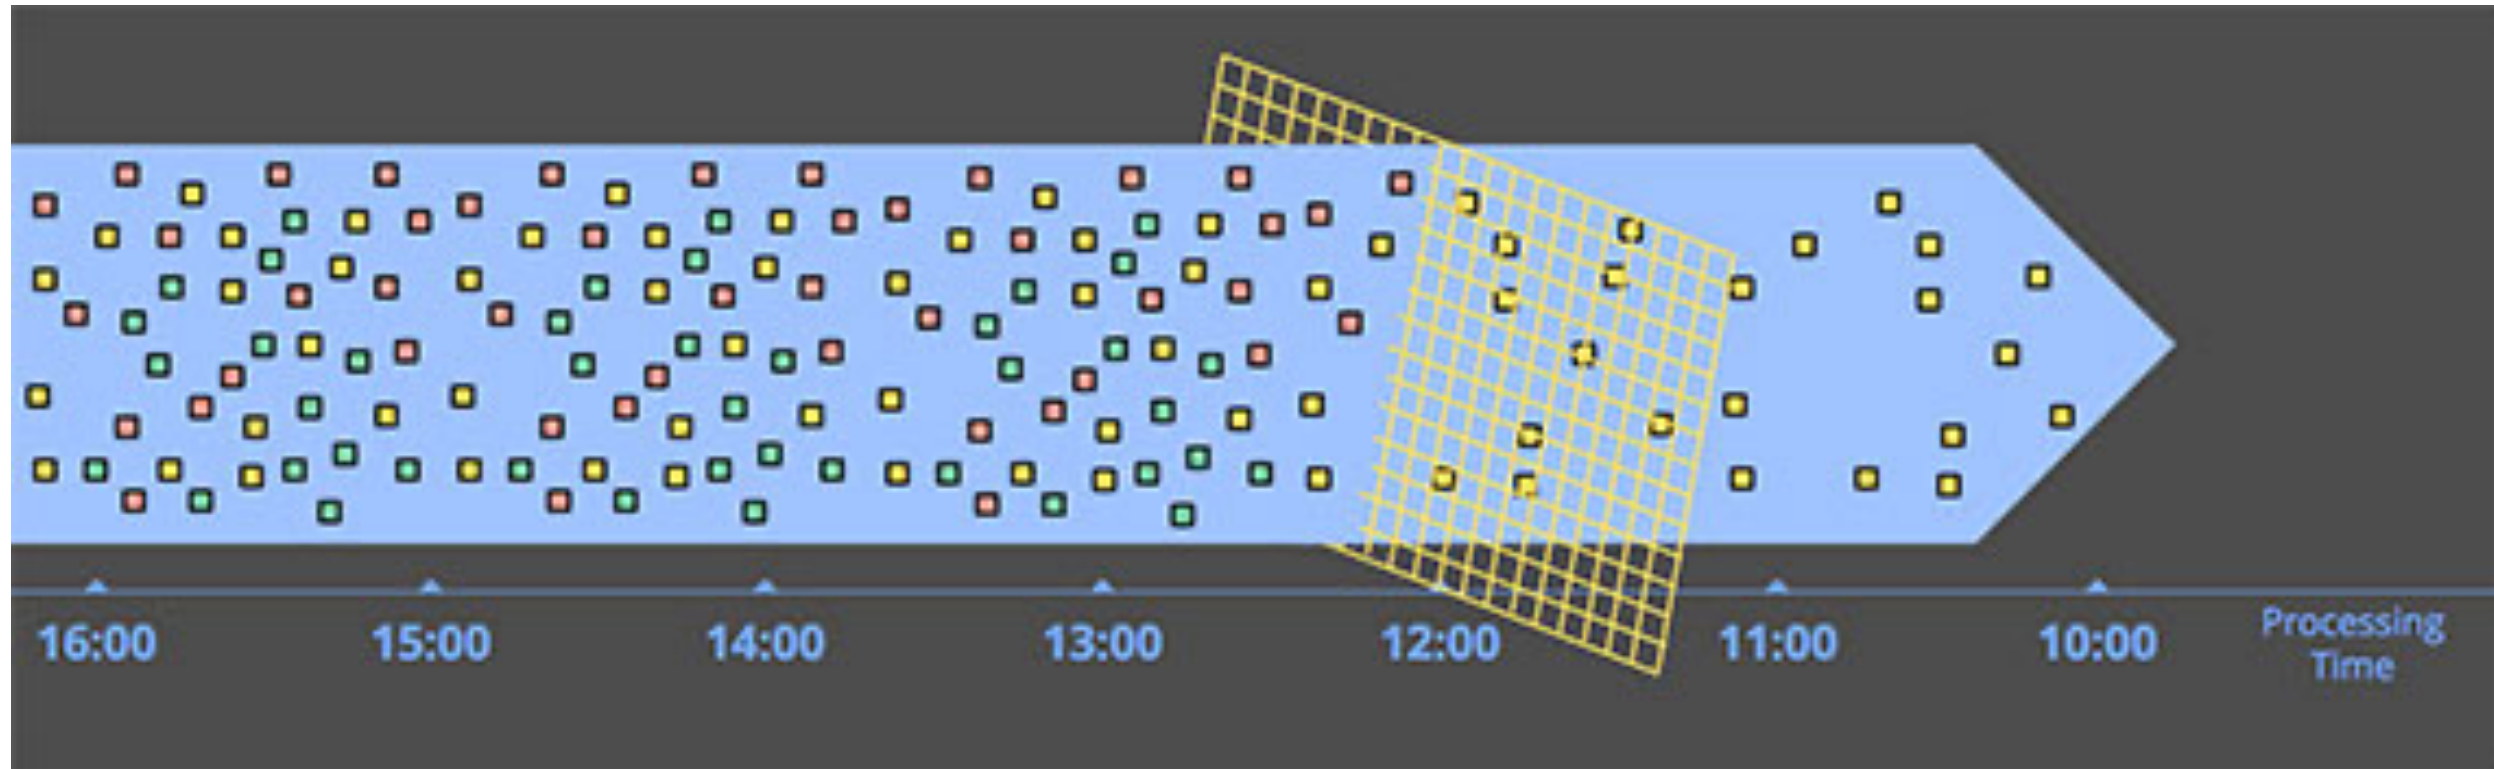


Agenda

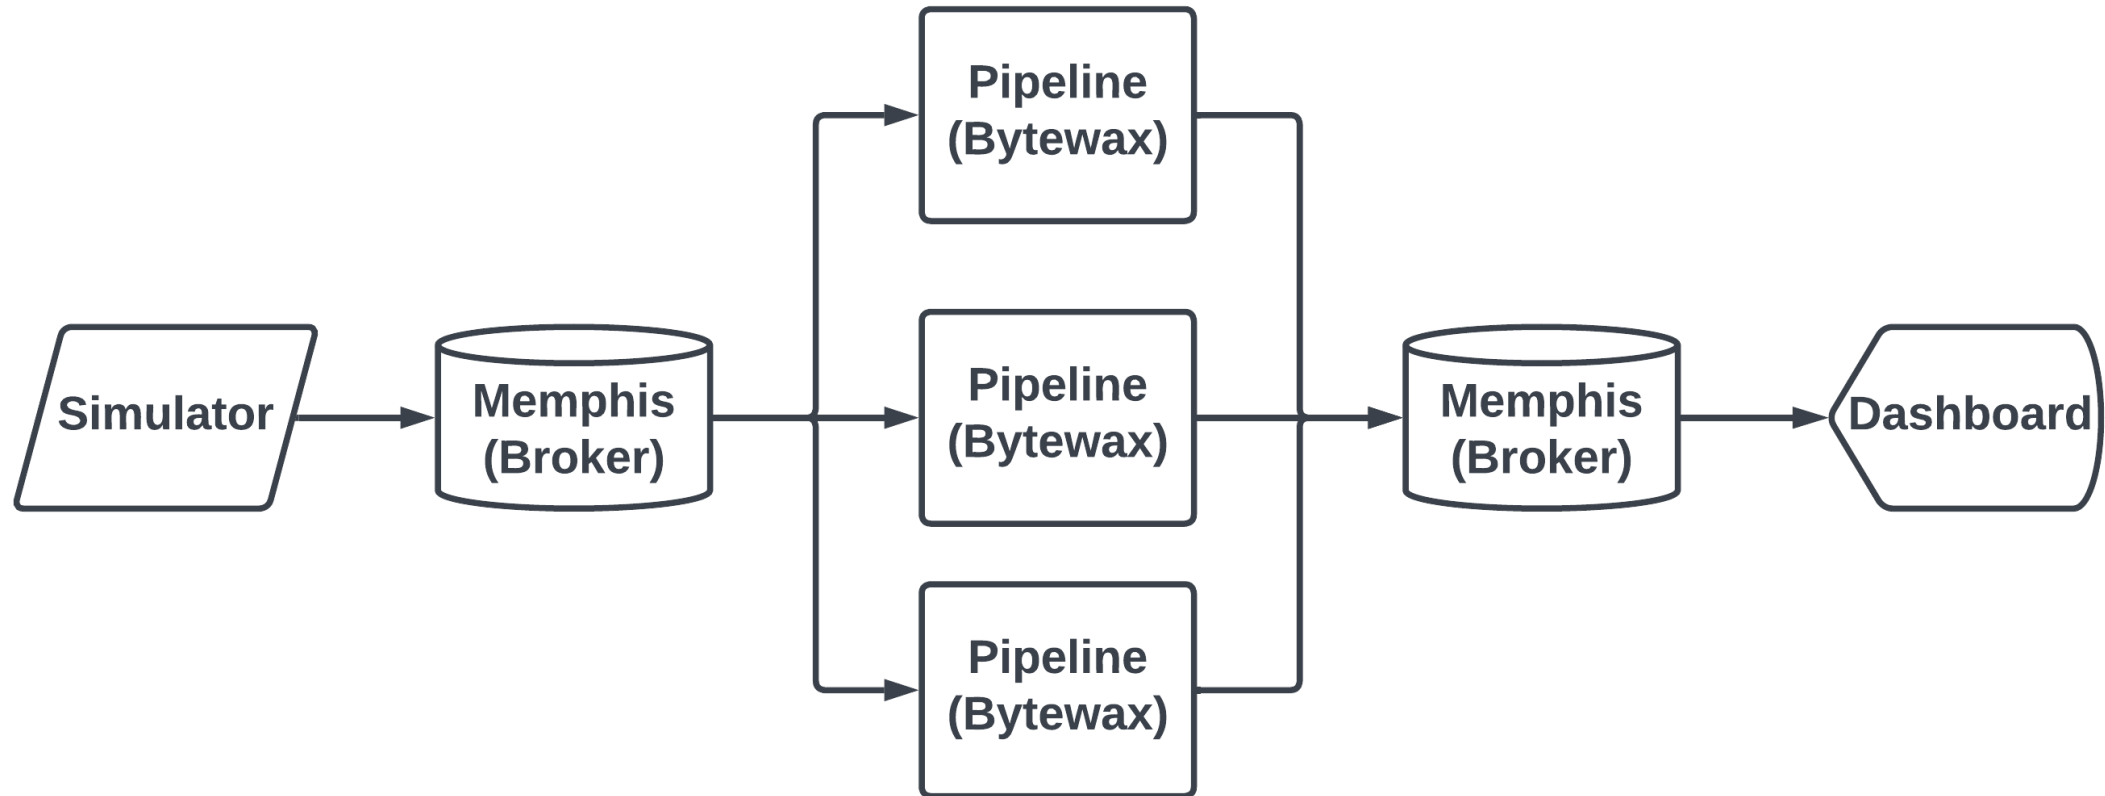
- Introduce message brokers
- Introduce stream processing engines
- Demo our transaction analysis system

Message Brokers

Streaming Data

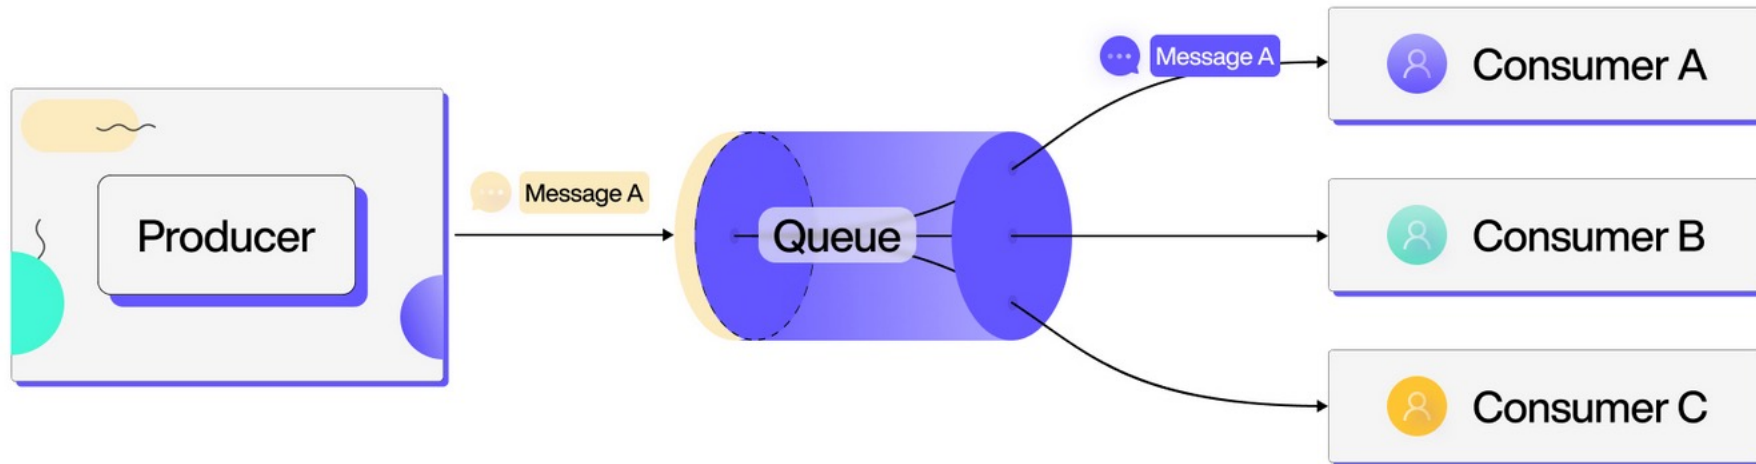


Architecture



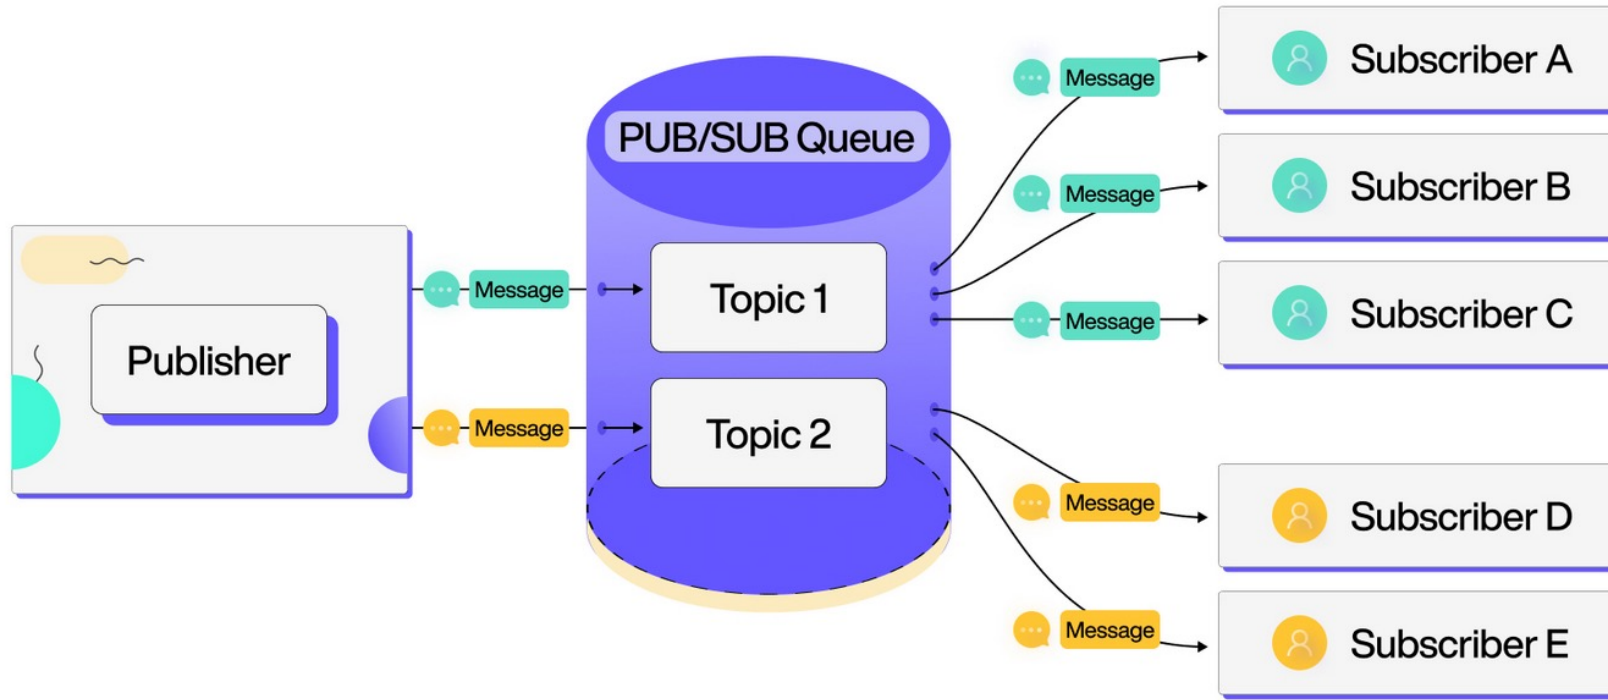
Message Queue Pattern

Point-to-point queues.

Publish / Subscribe Pattern

Publish-subscribe pattern.



Architectural Advantages

- Decouples routing
- Decouples availability
- Enables asynchronous communication
- Rate limiting / buffering

Producer Code

```
memphis = Memphis()
await memphis.connect(host=HOST, username=USERNAME,
                      password=PASSWORD)
producer = await memphis.producer(
    station_name=STATION,
    producer_name="test-producer")

msg_id = 0
while True:
    msg = f"This is test message {msg_id}."
    await producer.produce(bytearray(msg, "utf-8"))
```

Consumer Code

```
consumer = await memphis.consumer(  
    station_name=STATION,  
    consumer_name="test-consumer")  
  
while True:  
    batch = await consumer.fetch()  
    for msg in batch:  
        s = msg.get_data().decode("utf-8")  
        print(s)  
        msg.ack()
```

←

todo-cdc-events

+ Add new tag

Created by root at Apr 16, 2023, 5:55 PM

Retention: 7 days

Replicas: 1

Local Storage: disk

Remote Storage: —

</>

Schema

+ Attach schema

⌚

Total messages

8

📊

Av. message size

187 Bytes

Code examples

View details >

Audit

View details >

Purge

🗑

Producers (1)

Name	User	Status
test-producer	todocdcse...	🔴

Details

Name

test-producer

User

todocdcservice

IP

192.168.64.1:42380

Station

Showing last 8 out of 8 messages

Drop

Messages

Dead-letter

Details

Messages

Information

☐ This is a test me...

☐ This is a test me...

☐ This is a test me...

☐ This is a test me...

☐ This is a test me...

☐ This is a test me...

☐ This is a test me...

☐ This is a test me...

☐ This is a test me...

No message selected

Consumer groups (1)

Name	Unacked	Unprocessed	Status
test-consu...	0	0	🟢

Details

Unacked messages

0

Unprocessed messages

0

In process message

0

Max ack time

30,000ms

Max message deliveries

10

Consumers

1

>

"Classic" Message Brokers

- Apache ActiveMQ
- Apache RocketMQ
- BlazingMQ
- IBM MQ
- NATS
- Oracle Advanced Queuing
- RabbitMQ

"Modern" Message Brokers

- Apache Kafka
- Apache Pulsar
- Memphis
- NATS Jetstream
- Redpanda

Architectural Features

- Partitioning: enables parallelism
- Message retention: turns message broker into a database
- Consumer can rewind: message replay (time travel)

Stream Processing Engines

Stream Processing Engines

- Provide a programming model
 - "Automatic" parallelization
 - State management
 - Recovery
- Apache Flink
 - Apache Samza
 - Apache Spark
 - Apache Storm
 - Bytewax
 - Faust



```
flow = Dataflow()
flow.input("memphis-consumer", memphis_src)

# bytearray to UTF-8 string
flow.map(lambda m: m.decode("utf-8"))

# deserialize JSON document
flow.map(json.loads)

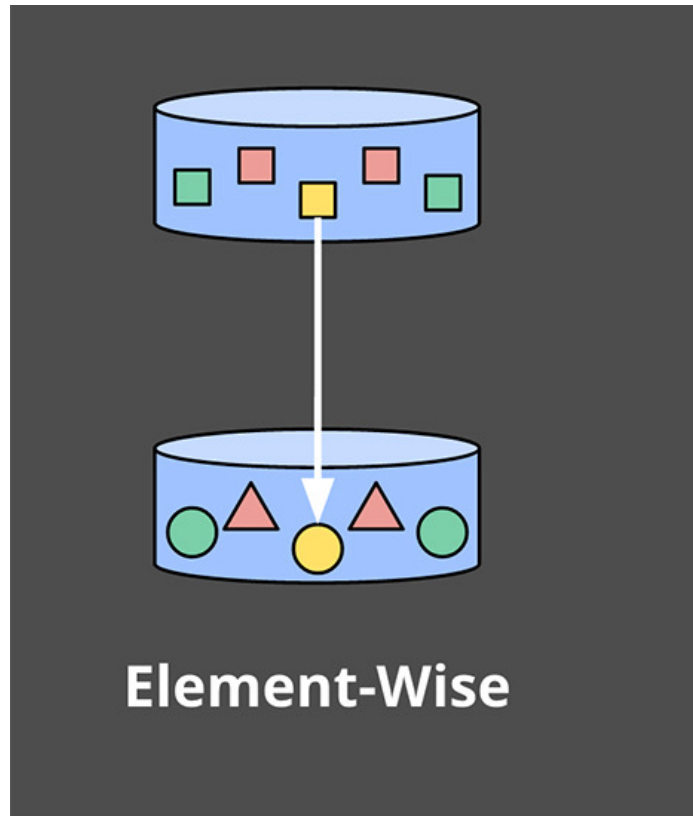
# grab customer id
flow.map(lambda r: r["customer_id"])

# convert to bytearray
flow.map(lambda s: bytearray(s, "utf-8"))

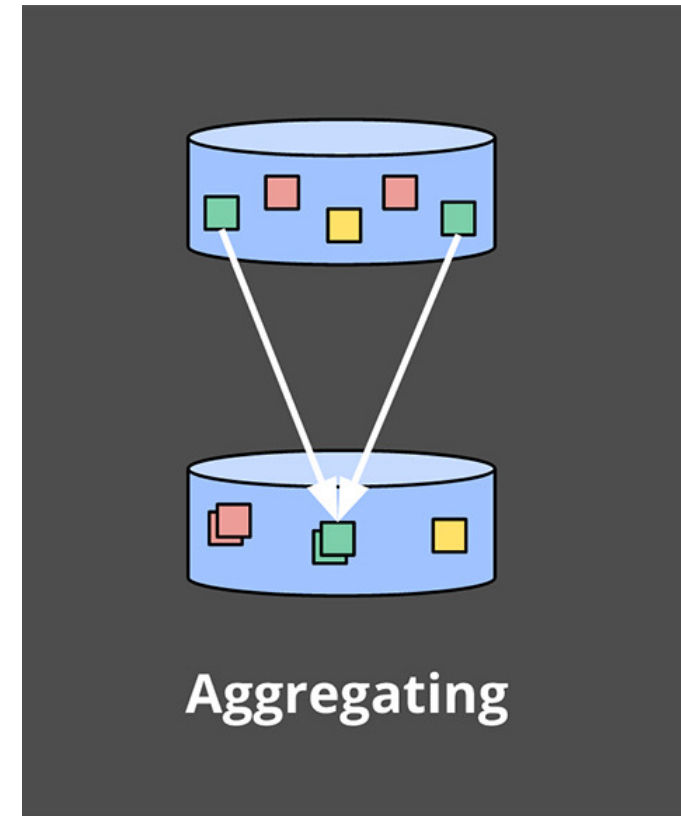
# print customer id
flow.output("console-output", StdOutput())
```

Transformations

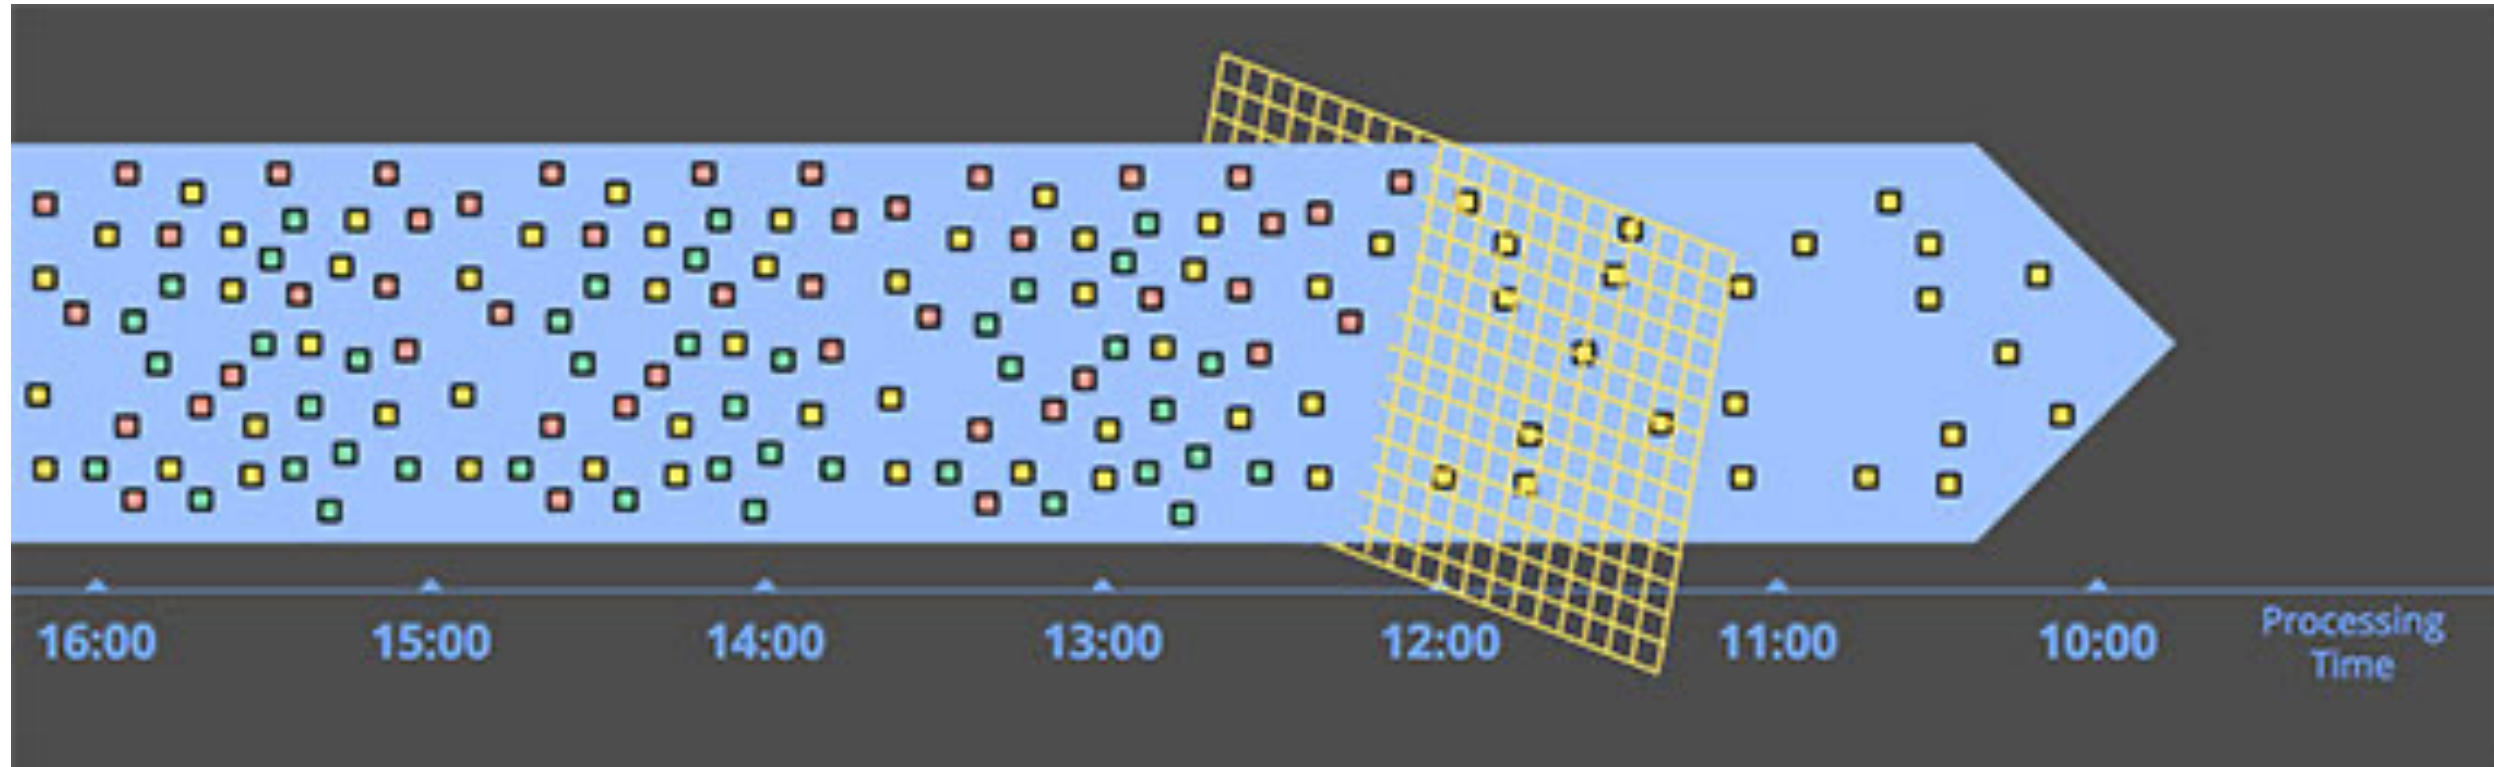
Map



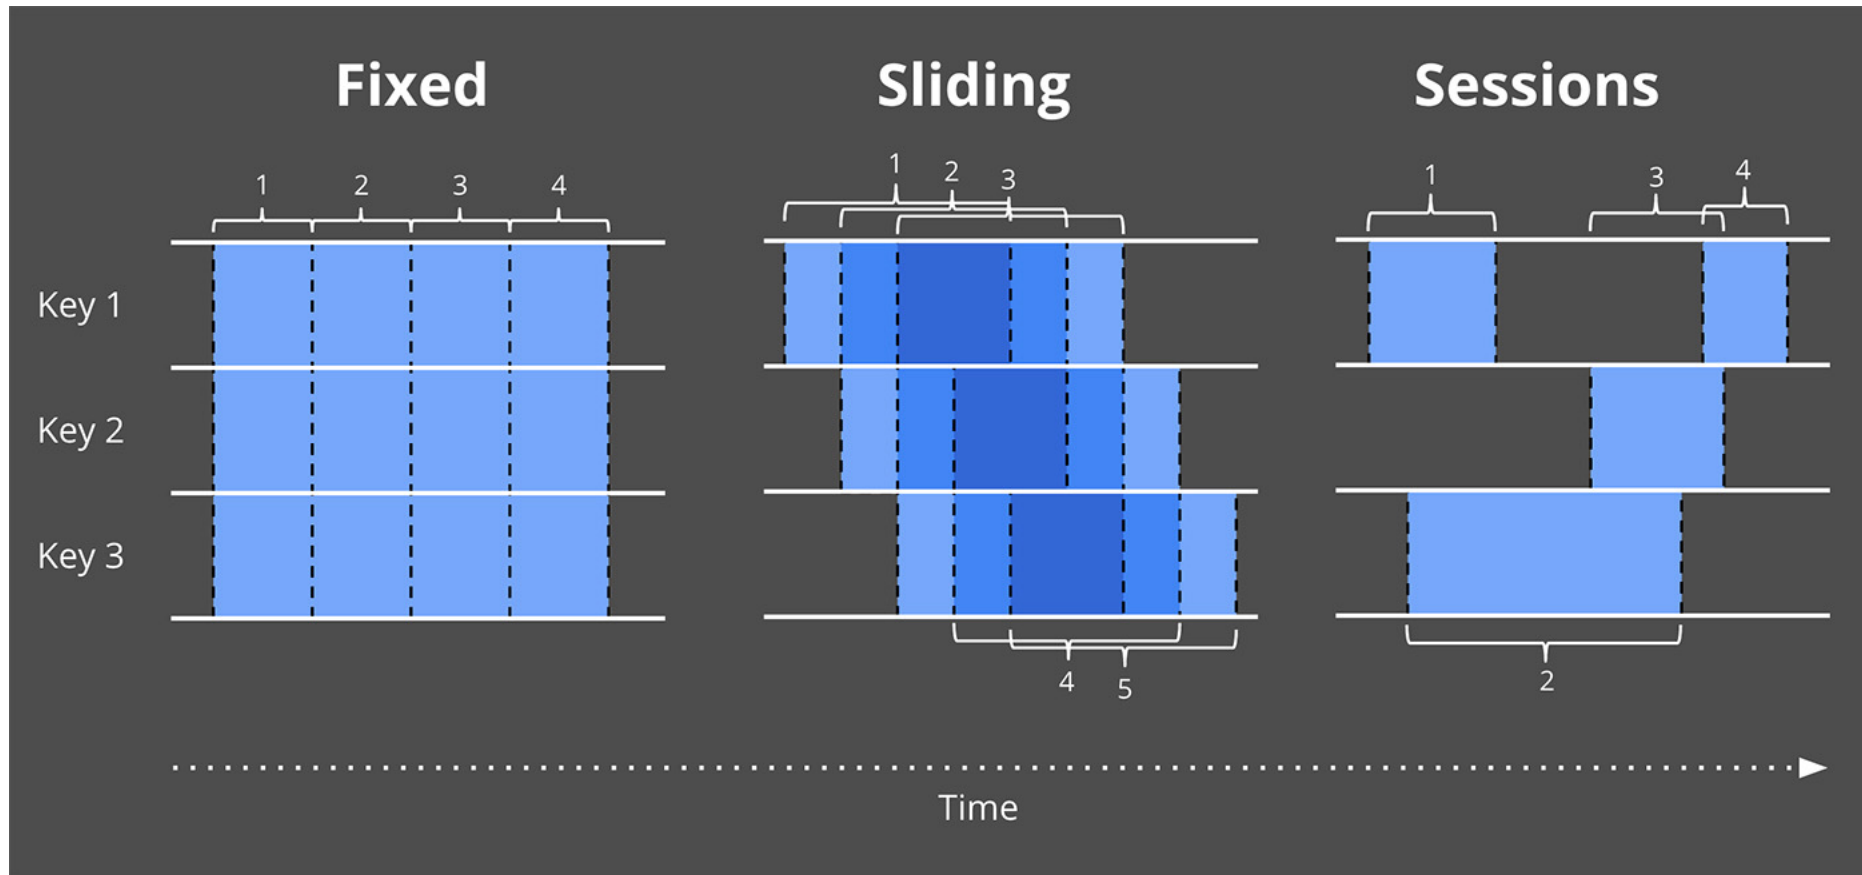
Reduce / Fold



Filter



Windowing



Demo

- Let's switch to our live demo...

memphis-example-solutionsPublic

Edit PinsWatch1Fork0Star1


master3 branches0 tags

Go to fileAdd fileCode


rnowling-memphis Remove curl (#18)5d5a39d 2 weeks ago28 commits

mongodb-debezium-cdc-example	Remove curl (#18)	2 weeks ago
postgres-debezium-cdc-example	Use Python Debian images instead of vanilla Debian images (#15)	last month
CODE_OF_CONDUCT.md	Add license, code of conduct, and README (#1)	3 months ago
LICENSE	Add license, code of conduct, and README (#1)	3 months ago
README.md	Add license, code of conduct, and README (#1)	3 months ago

README.md



Memphis is a next-generation alternative to traditional message brokers.



[Docs](#) - [Twitter](#) - [YouTube](#)

Introduction to this Repository

This repository contains examples of integrating [Memphis](#), the next-generation message broker, with other open-source technologies to solve technical use cases.

About

Example solutions using Memphis

- Readme
- View license
- Code of conduct
- Activity
- 1 star
- 1 watching
- 0 forks

Report repository

Releases

No releases published

Packages

No packages published

Contributors 3

- rnowling-memphis RJ Nowling
- rnowling RJ Nowling
- idanasulinmemphis Idan Asulin

Languages

Python 94.6% Dockerfile 5.4%



<https://github.com/memphisdev/memphis-example-solutions>

Thanks!